
Workshop 2:

GPU computing frameworks

Jeremy Jacobson
Lecturer



EMORY
COLLEGE
OF ARTS AND
SCIENCES

Department of Quantitative
Theory and Methods

Overview

Main points

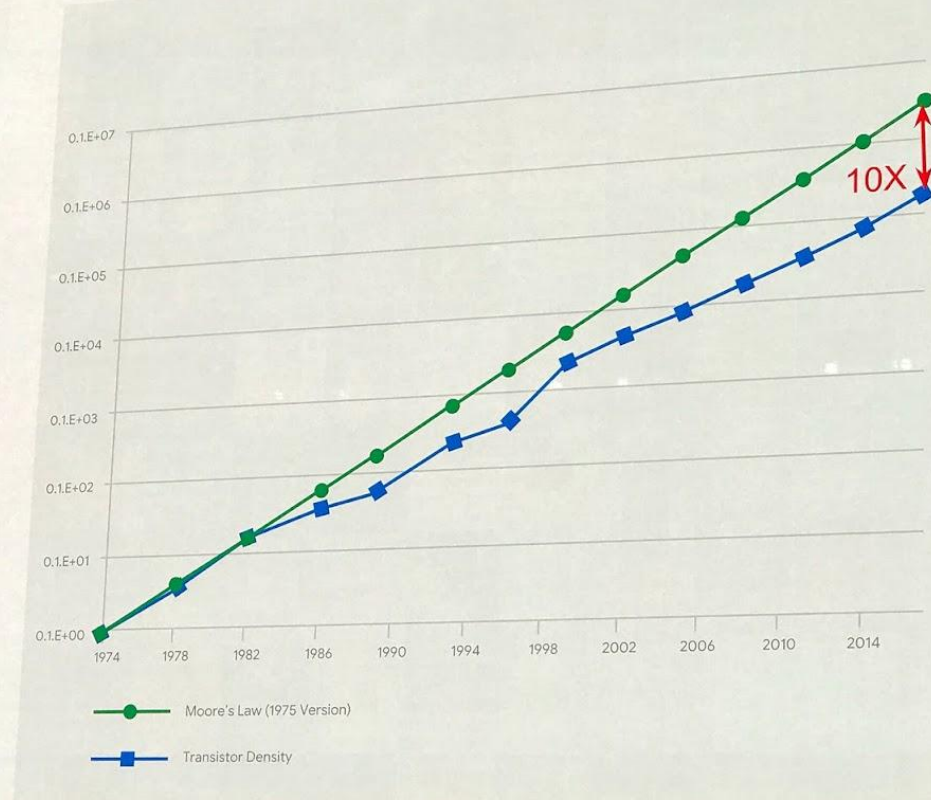
- (10 minutes) **PhD student Juan Estrada Sosa**
 - Postgres install and use via Docker on the QTM server
- (5 minutes) **Workshop 1 review**
- (30 minutes) **Python vrs Numpy vrs Pytorch**
 - We will revisit the matrix multiply speed improvements suggested in workshop 1 using these frameworks.

Workshop 1


review in three plots

Moore's Law slowdown in Intel Processors

transistor slowing down faster,
fab costs.



Deep learning is
causing a machine
learning revolution



From "A New Golden Age in Computer Architecture: Empowering the Machine-Learning Revolution."
Dean, J., Patterson, D., & Young, C. (2018). IEEE Micro, 38(2), 21-29.

What Opportunities Left?

- SW-centric
 - Modern scripting languages are interpreted, dynamically-typed and encourage reuse
 - Efficient for programmers but not for execution
- HW-centric
 - Only path left is Domain Specific Architectures
 - Just do a few tasks, but extremely well
- Combination
 - Domain Specific Languages & Architectures




A6000 (our GPU) exterior



What's the Opportunity?

Matrix Multiply: relative speedup
to a Python version (18 core Intel)

Matrix Multiply Speedup Over Native Python




from: "There's Plenty of Room at the Top," Leiserson, et. al., to appear.


Matrix multiply on a GPU

- Pure Python
- Numpy
- PyTorch

We continue by working through the workshop notebook

Each Tensor Core provides a 4x4x4 matrix processing array which performs the operation $\mathbf{D} = \mathbf{A} * \mathbf{B} + \mathbf{C}$





Graphics Card	NVIDIA RTX A6000	NVIDIA A40
GPU Codename	GA102	GA102
GPU Architecture	NVIDIA Ampere	NVIDIA Ampere
GPCs	7	7
TPCs	42	42
SMs	84	84
CUDA Cores / SM	128	128
CUDA Cores / GPU	10752	10752
Tensor Cores / SM	4 (3rd Gen)	4 (3rd Gen)
Tensor Cores / GPU	336 (3rd Gen)	336 (3rd Gen)
RT Cores	84 (2nd Gen)	84 (2nd Gen)
GPU Boost Clock (MHz)	1800	1740
Peak FP32 TFLOPS (non-Tensor) ¹	38.7	37.4
Peak FP16 TFLOPS (non-Tensor) ¹	38.7	37.4
Peak BF16 TFLOPS (non-Tensor) ¹	38.7	37.4
Peak INT32 TOPS (non-Tensor) ^{1,3}	19.4	18.7
Peak FP16 Tensor TFLOPS with FP16 Accumulate ¹	154.8/309.6 ²	149.7/299.4 ²
Peak FP16 Tensor TFLOPS with FP32 Accumulate ¹	154.8/309.6 ²	149.7/299.4 ²
Peak BF16 Tensor TFLOPS with FP32 Accumulate ¹	154.8/309.6 ²	149.7/299.4 ²
Peak TF32 Tensor TFLOPS ¹	77.4/154.8 ²	74.8/149.6 ²

Matrix multiply on a GPU

We continue by working through the
workshop notebook

References

- Slides from David Patterson's talk at Google Cloud Faculty Institute which was not recorded. He gave a talk on similar topics to the ACM and a recording is available [here](#).
 - [NVIDIA Ampere architecture whitepaper](#)
 - [CUDA programming guide](#)
-